



ASSIGNMENT 2

DESIGNING A DASHBOARD WITH REAL-WORLD DATA

Assignment Details:**Course Name:** Modern Business Intelligence**Course Code:** INTE2689**Lecturer:** Dr Hamed Vosgha**Submission Date:** 11 May 2026**Word Counts:** 2200**Prepared by:**

Ngo Hua Quoc Thinh - s3863887

Assessment Declaration:

I declare that in submitting all work for this assessment I have read, understood and agree to the content and expectations of the Assessment declaration.

Table of Contents

1. Introduction.....	3
2. Data Cleaning and Processing.....	3
3. Data Visualisation and Dashboard	7
3.1 Dashboard Design Overview	7
3.2 Key Measures and Calculated Fields.....	7
3.3 Dashboard 1: Geographic Performance.....	9
3.4 Dashboard 2: Nationality Bias Analysis	10
3.5 Dashboard 3: Sentiment Driver Analysis.....	11
3.6 Dashboard 4: Seasonality Performance	12
3.7 Business Value.....	13
4. TravelSphere's use cases	13
4.1 Use Case 1: Shortlisting Low-Risk Partner Hotels for a Barcelona Luxury Package	13
4.2 Use Case 2: Comparing Paris and Amsterdam for Higher-Expectation Nationality Segments.....	17
5. Summary of Business Value.....	23

1. Introduction

This BI dashboard is designed for TravelSphere, a luxury travel agency aiming to refine its European partner hotel portfolio using customer review data. The dashboard focuses on 2017 hotel reviews and supports the Head of Customer Experience in identifying high-performing locations, understanding reviewer nationality differences, detecting positive and negative sentiment drivers, and observing seasonal satisfaction patterns.

The dashboard uses a public hotel review dataset containing hotel names, addresses, review dates, reviewer nationalities, positive and negative review texts, reviewer scores, average hotel scores, word counts, and hotel coordinates. The purpose is not only to describe past customer feedback, but also to support partner hotel selection and client advisory decisions.

2. Data Cleaning and Processing

Key data preparation steps utilized *Tableau Prep Builder* in *Table 1*.

Table 1: Initial data cleaning steps in Tableau Prep Builder.

Step	Step's names	Cleaning / Processing Action	Rationale
<i>Step 1</i>	Date filtering	Converted <i>Review_Date</i> to date format and filtered <i>Range of Dates</i> : 01/01/2017 – 31/12/2017	Meets Head of Customer Experience's requirement
<i>Step 2</i>	Duplicate handling	<i>Identify Duplicate Rows</i> ; 'duplicate' records are excluded before final analysis	Avoids double-counting reviews
<i>Step 3</i>	Missing values	○ <i>lat</i> and <i>lng</i> have '<i>null</i>' values; these records were included in non-geographic analyses but excluded from map visualisations ○ '<i>null</i>' values of <i>Negative_Review</i> and <i>Positive_Review</i> will be handled in the <i>Step 7</i>	Prevents map errors, inaccurate location plotting while preserving reviews
<i>Step 4</i>	City extraction	○ Created <i>City</i> from <i>Hotel_Address</i>, including '<i>London</i>', '<i>Paris</i>', '<i>Barcelona</i>', '<i>Amsterdam</i>', '<i>Milan</i>', and '<i>Vienna</i>'. ○ Use <i>Create Calculated Field</i>: IF CONTAINS(UPPER([Hotel_Address]), " LONDON ") THEN "London" ELSEIF CONTAINS(UPPER([Hotel_Address]), " PARIS ") THEN "Paris"	support location-level comparison

		<pre>ELSEIF CONTAINS(UPPER([Hotel_Address]), " BARCELONA ") THEN "Barcelona" ELSEIF CONTAINS(UPPER([Hotel_Address]), " AMSTERDAM ") THEN "Amsterdam" ELSEIF CONTAINS(UPPER([Hotel_Address]), " MILAN ") THEN "Milan" ELSEIF CONTAINS(UPPER([Hotel_Address]), " VIENNA ") THEN "Vienna" ELSE "Others" END</pre> <i>Notes: City extraction was performed using a controlled city dictionary due to inconsistent address formats, especially UK addresses containing postcodes before the country name. In production, this should be replaced with a maintained location lookup table or geocoding service.</i>	
<i>Step 5</i>	Month transformation	<i>Review_Date</i> were transformed into two additional fields: <i>Review_Month</i> and <i>Review_Month_Number</i>. ○ <i>Review_Month</i> was created by <i>Create Calculated Field: DATENAME('month', [Review_Date])</i> ○ <i>Review_Month_Number</i> was created by <i>Create Calculated Field: DATEPART('month', [Review_Date])</i>	Support seasonality analysis by displaying monthly trends clearly while ensuring the months are sorted in the correct chronological order.
<i>Step 6</i>	1st Output	Export the 1st cleaned output into a .csv file. Name as ' <i>2017_Hotel_Reviews</i> '	To reduce file size and prepare a manageable cleaned dataset for AI-assisted text analysis.

After the initial cleaning, **ChatGPT** was used to analyse *Negative_Review* and *Positive_Review* text and extract further insights, including topics, reasons, and responsible teams. **Tableau Prep Builder** was then used to validate the AI-generated dataset and ensure it was ready for visualisation in **Tableau Desktop**. These steps are presented in the *Table 2* below.

Table 2: AI-assisted review text cleaning steps in Tableau Prep Builder.

<i>Step 7</i>	Text review standardising & Sentiment indicators	○ Upload <i>2017_Hotel_Reviews.csv</i> on ChatGPT, then input this prompt below: “Clean the 2017_Hotel_Reviews dataset and create a Tableau-ready CSV. Standardise neutral review values such as “No Negative”, “No Positive”, “Nothing”, “N/A”, “None”, “null”, and blanks so they are not misclassified as sentiment. Add	○ Extract root causes of satisfaction/dissatisfaction and ○ Improve file manageability for
---------------	--	--	---

		cleaned positive/negative review fields, adjusted word counts, sentiment balance/ratio/label, AI-assisted topic tags, satisfaction/dissatisfaction reasons, responsible team fields, and Tableau-friendly topic/team flag columns. Keep all original columns and export the final cleaned file as CSV.” ○ Then, input this prompt: “Seperate into smaller files” to split the large CSV into smaller files for download and re-upload on Tableau Prep	Tableau Prep processing manageability.
<i>Step 8</i>	Union file	Use ‘ <i>Union</i> ’ Function in Tableau Prep	Restore a single complete dataset with identifcal fields for Tableu Desktop
<i>Step 9</i>	Validate data type	‘Date’ format: <i>Review_Date</i> ‘Number’ format: <i>Reviewer_Score, Average_Score, Total_Number_of_Reviews</i> ‘Number/Boolean’ format: ‘<i>_Flag</i>’ fields. ‘String’ format: <i>Sentiment_Label, Responsible_Team, Dissatisfaction_Reason</i>	Ensure readability of data types in AI-generated fields
<i>Step 10</i>	Validate missing value	Check <i>Positive_Review_Cleaned</i> and <i>Negative_Review_Cleaned</i> that do not have ‘ <i>null</i> ’ values	Ensure fields do not contain unclear blank values
<i>Step 11</i>	Rename fields	Sentiment_Label → ‘<i>Review_Sentiment</i>’ Dissatisfaction_Reason → ‘<i>Reason_for_Dissatisfaction</i>’ Satisfaction_Reason → ‘<i>Reason_for_Dissatisfaction</i>’ Review_Topics_Combined → ‘<i>Review_Topic</i>’	Ensure readability of fields’ name
<i>Step 12</i>	Duplicate handling	<i>Identify Duplicate Rows</i> ; ‘duplicate’ records are excluded before final analysis	Avoids double-counting reviews
<i>Step 13</i>	Final Output	Export the final cleaned output into a <i>.csv</i> file. Name as ‘ <i>Final_2017_Hotel_Reviews</i> ’	Ready-to-visualised dataset for Tableau Desktop

The dataset was first filtered to include only reviews where *Review_Date* falls in 2017 (*Table 1*). After cleaning and AI-assisted text processing (*Table 2*), the final Tableau-ready dataset contains **156,681 review records**, **1,490 hotels**, and **206 reviewer nationalities**. The average reviewer score is approximately **8.39**, with scores ranging **from 2.5 to 10.0**, indicating generally high customer satisfaction. Although the dashboard focuses on 2017 reviews, the available review dates only cover January to early August 2017; therefore, seasonality analysis is limited to the available eight-month period and should not be interpreted as a full-year trend.

Several issues were identified during the cleaning process. First, some latitude and longitude values were missing, so these records were retained for general analysis but excluded from map visualisations to avoid inaccurate location plotting (*Table 1*). Second, hotel address formats were inconsistent, especially UK addresses, so a controlled city dictionary was used to extract cities reliably. Third, textual review fields contained misleading neutral values such as “*No Negative*”, “*No Positive*”, “*Nothing*”, “*N/A*”, “*None*”, “*null*”, and blank values. These were standardised to avoid misclassifying no complaint as negative sentiment. Main data concerns include incomplete coordinate fields, inconsistent review text formats, and potential nationality bias. As some nationalities have very low review counts, nationality comparisons should apply a minimum threshold, such as 100 or 500 reviews, to avoid misleading conclusions. Records with missing coordinates should be excluded from map visualisations but retained for non-geographic analysis. Finally, because the AI-enhanced CSV file was too large to download and process easily, it was split into smaller CSV files and later unioned in Tableau Prep. AI-assisted topic, reason, and responsible team fields were also treated as decision-support labels rather than perfect manual classifications.

The Tableau Prep Builder workflow is included in *Figure 1* to visually summarise the data cleaning process described in *Table 1* and *Table 2*. It facilitates how the raw hotel review data was transformed through initial cleaning, AI-assisted text analysis, validation, union of split files, and final output for Tableau Desktop visualisation.



Figure 1: Tableau Prep Builder workflow for Hotel_Reviews data preparation.

3. Data Visualisation and Dashboard

3.1 Dashboard Design Overview

The Tableau dashboard report was developed for TravelSphere’s Head of Customer Experience to support data-driven decisions in partner hotel selection and luxury client advisory. The final report includes four interactive dashboards: ***Dashboard (1): Geographic Performance***, ***Dashboard (2): Nationality Bias Analysis***, ***Dashboard (3): Sentiment Driver Analysis***, ***Dashboard (4): Seasonality Performance***. Together, these dashboards translate the cleaned 2017 hotel review dataset into an executive-level decision-support tool, addressing key business questions related to location performance, reviewer nationality, sentiment patterns, and monthly trends.

The dashboard design is based on the cleaned and AI-enhanced dataset, including review scores, cities, reviewer nationalities, sentiment topics, responsible teams, and monthly fields. A consistent executive layout was applied across all dashboards, using a blue header, navigation tabs, executive insight box, interactive filters, and concise chart titles. This structure supports quick interpretation and follows a top-down analytical flow: from high-level KPIs, to diagnostic charts, to operational details. As a result, the Head of Customer Experience can move from understanding “*what is happening*”, to identifying “*why it is happening*”, and finally deciding “*what action should be taken*”.

3.2 Key Measures and Calculated Fields

Several key measures and calculated fields were created to support the dashboard analysis. *Average Reviewer Score* was used as the main satisfaction measure to compare cities, hotels, nationalities, and monthly performance. *Review Count* was included as a reliability indicator, supported by minimum review threshold filters to reduce the risk of interpreting results from very small sample sizes.

A *Satisfaction Band* was created to classify hotels or cities into high, moderate, and low satisfaction groups, making score differences easier for management to interpret. In the partner hotel table, a *Risk Level* measure was also used to assess hotels based on both satisfaction performance and negative comment percentage. This supports portfolio decisions by identifying hotels that may appear highly rated but still require further investigation due to recurring negative feedback.

For sentiment analysis, AI-assisted fields were used to classify *Satisfaction Topic*, *Dissatisfaction Topic*, *Responsible Teams*, and *Recommended Team Action*. These fields convert unstructured review text into decision-ready categories such as room, staff, breakfast, facilities, location, cleanliness, and value. Since these AI-generated labels are intended as decision-support indicators, the dashboard focuses on identifying patterns and priorities rather than claiming absolute causality.

For seasonality analysis, *Review Month* and *Review Month Number* were created to display monthly trends in chronological order. These fields support the analysis of review volume, satisfaction score, and monthly negative drivers across the available 2017 review period.

3.3 Dashboard Design Rationale and Interactivity

The dashboard design was developed to support fast executive interpretation and interactive exploration. A consistent blue header was used across all dashboards to create a professional visual identity for TravelSphere and maintain design consistency. KPI cards were placed at the top of the ***Dashboard (1)*** to provide an immediate summary of total reviews, total hotels, average score, and high satisfaction percentage before users explore detailed charts.

Different chart types were selected based on their analytical purpose. Bar charts were used for ranking cities, hotels, and nationalities because they allow direct comparison. The map was used to show geographic hotel performance and identify location-based satisfaction clusters. Scatter plots were used in the ***Dashboard (2)*** to compare nationality review volume and average score, helping users identify both rating behaviour and sample reliability. Heatmaps were used in the ***Dashboard (3)*** to reveal satisfaction patterns across city and month, while tables were used where detailed operational decision-making was required.

Colour was applied consistently to support interpretation. Blue tones were used for positive or high-satisfaction results, while red and orange tones were used to highlight dissatisfaction, negative sentiment, or higher-risk areas. This colour logic allows the Head of Customer

Experience to identify strong performance and service risk quickly. Threshold filters were also included to improve analytical reliability by allowing users to remove hotels or nationalities with very low review counts.

The dashboards include interactive filters for city, hotel name, review month, satisfaction band, reviewer score band, dissatisfaction topic, and minimum review thresholds. These filters allow users to move from a broad executive overview to a focused diagnostic view. For example, the Head of Customer Experience can select one city in the ***Dashboard (1)***, then use the ***Dashboard (4)*** to understand the main causes of dissatisfaction in that market. This interaction supports a digital storytelling flow from market performance, to customer segment behaviour, to operational improvement priorities.

3.3 Dashboard 1: Geographic Performance




Figure 2: Dashboard 1 - Geographic Performance in Tableau Desktop.

The ***Dashboard (1)*** helps the Head of Customer Experience identify high-performing cities and hotel locations. It includes four *KPI cards*: total reviews, total hotels, average reviewer score, and high satisfaction percentage, providing a quick overview of dataset coverage and overall performance (*Figure 2*).

The dashboard combines a *City Satisfaction Ranking*, *Hotel Performance Map*, and *Top Partner Hotel Candidates* table (*Figure 2*). Barcelona and Vienna show stronger satisfaction

performance, suggesting stronger partnership potential in these markets. The hotel table adds operational detail by showing review count, average score, negative comment percentage, and risk level (*Figure 2*), helping management shortlist reliable partner hotels and identify properties requiring further investigation.

Interactive filters for *city*, *review month*, *satisfaction band*, and parameter *minimum hotel reviews* (*Figure 2*) allow users to focus on statistically reliable results and compare hotel performance across markets.

3.4 Dashboard 2: Nationality Bias Analysis




Figure 3: Dashboard 2 - Nationality Bias Analysis in Tableau Desktop.

The ***Dashboard (2)*** examines whether reviewer nationalities rate hotels differently. This is important because TravelSphere’s client recommendations may need to consider cultural expectations and rating behaviour.

The ***Dashboard (2)*** uses a *Nationality Bias Scatter Plot* to compare nationality review volume and average score, supported by a reference line showing the overall average (*Figure 3*). The *Harsh vs Generous Reviewer Nationalities* chart further shows each nationality’s score gap from the overall average, while the *lowest* and *highest rating nationality* charts provide quick comparison (*Figure 3*).

A *nationality review threshold* parameter was added to reduce misleading conclusions from small sample sizes (Figure 3). Additional filters for city, hotel name, and review month allow the Head of Customer Experience to examine whether nationality bias varies by location, property, or time period.

3.5 Dashboard 3: Sentiment Driver Analysis




Figure 4: Dashboard 3 - Sentiment Driver Analysis in Tableau Desktop.

The **Dashboard (3)** explains the main reasons behind customer satisfaction and dissatisfaction. While score-based dashboards show performance levels, this dashboard identifies the root causes behind those scores.

The *Positive* and *Negative sentiment* charts summarise key review topics such as staff, room, location, breakfast, facilities, cleanliness, and value (Figure 4). The *Responsible Team Priorities* table links negative review themes to internal teams, turning customer feedback into operational accountability (Figure 4). The *Positive vs Negative Topic Matrix* shows how satisfaction and dissatisfaction topics co-occur within the same reviews (Figure 4), helping identify mixed customer experiences.

The *Top 10 Negative Sentiments* table provides the most actionable view by showing review topic, responsible team, recommended action, review count, and percentage of negative reviews (Figure 4). This helps TravelSphere prioritise improvement areas before

recommending hotels to luxury clients. Filters for city, hotel name, review month, and reviewer score band support targeted diagnosis.

3.6 Dashboard 4: Seasonality Performance




Figure 5: Dashboard 4 - Seasonality Performance in Tableau Desktop.

The **Dashboard (4)** examines how review volume, satisfaction, and negative drivers change across the available 2017 months. This is important because hotel performance may fluctuate during peak travel periods due to higher demand and operational pressure.

The *Monthly Review Volume and Satisfaction* chart compares review count with average reviewer score, showing whether higher demand coincides with lower satisfaction (*Figure 5*). For example, July records high review activity while also showing lower average satisfaction, suggesting possible peak-season service strain. The *Monthly Negative Drivers* chart identifies which complaint topics increase by month, while the *City by Month Satisfaction Heatmap* compares satisfaction patterns across cities and months (*Figure 5*).

Filters for city, hotel name, review month, *dissatisfaction topic*, and parameter minimum hotel reviews allow targeted seasonality diagnosis. Since the available 2017 data only covers January to early August, findings should be interpreted as partial-year trends rather than a complete annual pattern.

3.7 Business Value

The Tableau dashboard report helps the Head of Customer Experience make evidence-based decisions in three areas. First, it supports partner hotel selection by identifying high-performing cities and reliable hotels with strong satisfaction, sufficient review volume, and low risk. Second, it improves customer advisory by showing how ratings vary across reviewer nationalities and travel periods. Third, it supports service improvement by linking negative review topics to responsible teams and recommended actions.

Therefore, the dashboards do not only describe historical review data but also provide a practical decision-making framework for TravelSphere. The report creates a digital narrative from market-level opportunity to hotel-level selection, to operational improvement, helping TravelSphere refine its European hotel portfolio and deliver more personalised luxury travel experiences.

4. TravelSphere's use cases

4.1 Use Case 1: Shortlisting Low-Risk Partner Hotels for a Barcelona Luxury Package

This use case demonstrates how TravelSphere can use the dashboards to select reliable partner hotels for a luxury travel package in Barcelona. The Head of Customer Experience begins with the ***Dashboard (1)*** and compares city-level satisfaction using the ***City Satisfaction Ranking (Figure 2)***. Barcelona is identified as one of the stronger-performing cities (Average Score = 8.56 with Minimum Hotel Review = 100), making it a potential priority market for partner hotel selection (*Figure 6*).




Figure 6: Geographic Performance dashboard filtered to Barcelona.

The Head then applies the *City* filter to Barcelona and increases the *Minimum Hotel Reviews* threshold (at 200 minimum reviews) (Figure 7). This step is important because a hotel with a very high score but only a small number of reviews may not be reliable enough for luxury client recommendations. The *Top Partner Hotel Candidates* table is then used to compare hotels based on average reviewer score, review count, negative comment percentage, and risk level (Figure 7). Instead of selecting the hotel with the highest average score only, TravelSphere should prioritise hotels that meet three conditions: strong average score, sufficient review volume, and low negative comment percentage, in this case is ‘*Avenida Palace*’ (Figure 7). A hotel with a high score but a high negative comment percentage should be treated as a warning sign, because luxury clients may be more sensitive to recurring service issues.




Figure 7: Top Partner Hotel Candidates table with minimum review threshold applied.

After shortlisting a candidate hotel, the user moves to the ***Dashboard (3)*** and filters by ‘Avenida Palace’ and the 4 nearest review months (Figure 8). This step explains the reasons behind the hotel’s performance. Since negative sentiment is mainly related to room quality (Negative Sentiment – Room = 44.18% of total) (Figure 8), this is a serious concern because room experience is central to luxury travel. In this case, TravelSphere should request quality assurance from the hotel before including it in the luxury package.




Figure 8: Sentiment Driver Analysis dashboard filtered by Barcelona, selected ‘Avenida Palace’, and the 4 nearest months.

The Head of Customer Experience then checks the ***Dashboards (4)*** to assess whether Barcelona’s satisfaction remains stable during high-demand months (*Figure 9*). However, July shows high review volume (7.89) but lower satisfaction in *Room* quality (198 review count), this suggests potential summer service pressure (*Figure 9*). In conclusion, TravelSphere still might recommend ‘Avenida Palace’ in Barcelona since stable monthly satisfaction but can consider other hotels with stable satisfaction and lower negative sentiment during peak months with higher minimum hotel reviews.

Overall, this use case shows how the dashboard supports a more reliable partner selection process. The final decision is not based only on average rating, but on a combination of city performance, hotel-level reliability, risk level, sentiment root causes, and seasonal stability.




Figure 9: Seasonality Performance dashboard showing monthly satisfaction and review volume.

4.2 Use Case 2: Comparing Paris and Amsterdam for Higher-Expectation Nationality Segments

This use case demonstrates how TravelSphere can use **Dashboard (2)** to personalise hotel recommendations for different customer segments (Figure 3). The Head of Customer Experience first identifies nationalities that rate below the overall average, as lower scores may reflect stricter expectations rather than poor hotel quality alone.

The *Nationality Bias Scatter Plot* and *Harsh vs Generous Reviewer Nationalities* compares average score, review volume, and score gap from the overall average by nationality (Figure 10). Nationalities below the overall average reference line of 8.23, including *India* (-0.57), *Saudi Arabia* (-0.53), and the *United Arab Emirates* (-0.52), can be interpreted as stricter reviewer groups (Figure 10). For TravelSphere, these groups should be treated as higher-expectation customer segments. It helps the business recognise where more careful hotel matching, clearer pre-trip communication, or stronger service standards may be required. To improve reliability, the *Nationality Review Threshold* is set to 200 reviews, reducing the risk of drawing conclusions from very small samples (Figure 10).




Figure 10: Nationality Bias Analysis dashboard showing nationality score gaps and Nationality Bias Scatter Plot and review threshold applied.

The user then filters **Dashboard (2)** by city to compare whether stricter nationality groups show the same rating behaviour in 'Paris' and 'Amsterdam'. In Amsterdam, *India* (-0.64) and the *United Arab Emirates* (-0.52) remain below the city average, suggesting that these segments may still have higher expectations in this market (*Figure 11*). In Paris, the pattern becomes stronger, with *India* (-0.80), *Saudi Arabia* appearing below average, and the *United Arab Emirates* (-0.66) also rating lower (*Figure 12*). This indicates that Paris may create a larger expectation gap for these customer groups than Amsterdam. Therefore, TravelSphere should apply stricter hotel screening when recommending Paris hotels to clients from India, Saudi Arabia, and the UAE, especially by checking whether dissatisfaction is linked to rooms, staff service, or facilities.




Figure 11: Nationality Bias Analysis dashboard showing nationality score gaps in Amsterdam.




Figure 12: Nationality Bias Analysis dashboard showing nationality score gaps in Paris.

To understand the reasons behind lower ratings, the user moves to **Dashboard (3)** compares Amsterdam and Paris. In both cities, dissatisfaction is mainly driven by *Room-related* issues, followed by *Staff* (Figure 13 & 14). However, Paris shows a slightly higher concentration of

room-related complaints, with Room accounting for 46.87% of negative sentiment compared with 43.73% in Amsterdam. Paris also shows stronger combined issues around *Room; Location* (19.54%) and *Room; Facilities* (15.00%) (Figure 13), while Amsterdam's key issues are *Room* (29.73%), *Room; Facilities* (14.79%), and *Room; Location* (11.71%) (Figure 14). This suggests that Paris presents a higher physical product risk for stricter nationality segments, especially around room quality, location expectations, and facilities. Therefore, TravelSphere should apply stricter hotel screening in Paris, while Amsterdam may be comparatively safer but still requires checks for recurring room and facilities complaints.




Figure 13: Sentiment Driver Analysis dashboard showing dissatisfaction issues in Paris.




Figure 14: Sentiment Driver Analysis dashboard showing dissatisfaction issues in Amsterdam.

Finally, the user checks ***Dashboard (4)*** to compare whether Amsterdam and Paris become riskier during busy months. In Amsterdam, review volume peaks in July (1,520 reviews) while satisfaction remains relatively stable at 7.86, suggesting manageable seasonal pressure (*Figure 15*). In Paris, July also has the highest review volume (314 reviews), but satisfaction remains lower overall at 7.57, following weaker scores in May and June (*Figure 16*). This suggests Paris may carry higher seasonal risk for stricter nationality segments. Therefore, TravelSphere should be more cautious when recommending Paris during peak months, prioritise higher-rated hotels, and provide clearer pre-trip expectation-setting; Amsterdam appears comparatively more stable but should still be monitored for July demand pressure.




Figure 15: Seasonality Performance dashboard showing review volume peak in Amsterdam.




Figure 16: Seasonality Performance dashboard showing review volume peak in Paris.

Overall, the dashboard suggests that Paris carries higher service risk than Amsterdam for stricter nationality segments, particularly clients from India, Saudi Arabia, and the United Arab Emirates. While both cities show recurring room-related dissatisfaction, Paris has lower overall satisfaction and stronger complaints around room quality, location, and facilities, especially

during higher-demand months. Therefore, TravelSphere should apply stricter hotel screening when recommending Paris, prioritising hotels with strong room quality, reliable staff service, and limited facilities complaints. Amsterdam appears comparatively more stable, but hotels should still be checked for recurring room and facilities issues before being recommended to luxury clients.

5. Summary of Business Value

The developed Tableau dashboards provide TravelSphere with a practical executive tool for improving hotel partner selection and customer advisory. The dashboards are not limited to descriptive reporting; they create an interactive decision-making process. The Head of Customer Experience can identify high-performing cities, shortlist reliable hotels, understand nationality-based rating differences, diagnose service issues, and monitor seasonal satisfaction patterns. As a result, the dashboard supports TravelSphere's strategic goal of refining its European hotel portfolio and delivering more personalised luxury travel experiences.